

CS6780

Summary and Conclusions

CS6780 – Advanced Machine Learning
Spring 2019

Thorsten Joachims
Cornell University

Theme: Prediction and Action

- Building intelligent systems vs. analyzing existing systems
 - Prediction
 - Intelligent action
 - Guarantees on prediction/action quality
 - Causality

Theme: Bias vs. Variance

- Fundamental trade-off in learning
 - Training error vs. prediction error
 - Model capacity
 - Statistical learning theory
 - Empirical risk minimization

Theme: Massive Overparameterization

- The success story of machine learning
 - Regularized linear models
 - Kernels
 - Deep networks
 - Number of parameters \gg number of examples

Theme: Theoretical Underpinning

- Theory for understanding sake
 - Identify the mechanisms at play in ML
 - Understand model complexity
 - Understand common themes between algorithms

Secondary Syllabus

- Practice “soft skills” needed to be a successful researcher
 - Pitch ideas
 - Present your work
 - Write convincing papers
 - Work in groups
 - Give constructive feedback to others
 - Use feedback constructively

Batch Learning Approaches

- Empirical Risk Minimization (ERM)
 - Fixed at training time: class of decision rules $h: X \rightarrow Y$, loss, x and y
 - Strategy: minimize training loss
- Conditional Probability Models
 - Fixed at training time: class of models for $P(Y|X)$, x and y
 - Strategy: max conditional likelihood or MAP (or Bayes)
- Generative Models
 - Fixed at training time: class models for $P(Y,X)$
 - Strategy: max likelihood or MAP (or Bayes)

Not covered: Bayesian ML perspective
→ ORIE 6741, Andrew Wilson

Batch Learning for Classification

- ERM
 - Decision Trees
 - Perceptron
 - Linear SVMs
 - Kernel SVMs
 - Neural Networks
- Conditional Probability
 - Logistic Regression
 - Conditional Random Fields
 - Ridge Regression
- Generative
 - Multinomial Naïve Bayes
 - Multivariate Naïve Bayes
 - Linear Discriminant
- Other Methods
 - Logical rule learning
 - Gaussian Processes
 - RBF Networks
 - Boosting
 - Bagging
 - Parametric (Graphical) Models
 - *-Regression
 - *-Multiclass

→ Methods + Theory + Algorithms + Practice

Batch Learning for Struct Prediction

- ERM
 - Structural SVMs
- Conditional Probability
 - Conditional Random Fields
- Generative
 - Hidden Markov Model
- Other Methods
 - Maximum Margin Markov Networks
 - Markov Random Fields
 - Bayesian Networks
 - Statistical Relational Learning
 - Markov Logic Networks
 - Encoder/Decoder Networks

NLP classes

→ Claire Cardie

Online Learning

- Expert Setting
 - Halving Algorithm
 - Weighted Majority
 - Exponentiated Gradient
- Bandit Setting
 - EXP3
 - UCB1
- Other Methods
 - Hedge
 - Follow the Leader
 - Zooming
 - Partial Monitoring
 - Contextual Bandits
 - Dueling Bandits
 - Coactive Learning

Unsupervised Learning

- Clustering
 - K-Means
 - Mixture of Gaussians and EM-Algorithm
- Other Methods
 - Spectral Clustering
 - Multi-Dimensional Scaling
 - Latent Dirichlet Allocation
 - Semantic Embeddings
 - Deep Auto-Encoders
- Other Tasks
 - Outlier Detection
 - Novelty Detection
 - Dimensionality Reduction
 - Non-Linear Manifold Detection
- CS4786 Machine Learning for Data Science
- CS4850 Math Found for the Information Age
- INFO 6150 Advanced Topic Models

ML in Computer Visions

- Covered
 - Convolutional Neural Networks (LeNet)
- Other
 - More Deep Learning
 - Even more Deep Learning
 - CS6670 Computer Vision
 - Bharath Hariharan
 - CS4670 Intro Computer Vision
 - Serge Belongie
 - Bharath Hariharan

Learning to Act

- Covered
 - Off-policy policy learning
 - Atomic actions
 - Combinatorial actions
- Other
 - Sequential decisions
 - Markov Decision Processes
 - Reinforcement learning
 - On policy vs. off policy
 - Policy gradient

ML and Causality

- Covered
 - Potential outcomes model
 - Treatment effect estimation
- Other
 - Observational setting
 - Instrumental variables
 - Continuous treatments
 - Longitudinal treatments
 - Causal discovery
 - Parameter inference
 - Causal networks
 - Structural equation models

ML and Fairness

- Covered
- Other
 - Fairness
 - Accountability
 - Transparency
 - Criteria and policy
 - Algorithms and guarantees
 - Classification
 - Ranking

INFO4270: Ethics and Policy in DS

→ Solon Barocas

Other Machine Learning Courses at Cornell

- INFO 3300 - Data-Driven Web Pages
- CS 4700 - Introduction to Artificial Intelligence
- CS 4780/5780 - Machine Learning for Intelligent Systems
- CS 4786/5786 - Machine Learning for Data Science
- CS 4787 - Principles of Large-Scale Machine Learning
- OR 4740 - Statistical Data Mining
- CS 6780 - Advanced Machine Learning
- CS 6783 - Machine Learning Theory
- CS 6784 - Advanced Topics in Machine Learning
- CS 6787 - Advanced Machine Learning Systems
- ORIE 6740 - Statistical Learning Theory for Data Mining
- ORIE 6741 - Bayesian Machine Learning
- ORIE 6750 - Optimal learning
- ORIE 6780 - Bayesian Statistics and Data Analysis
- INFO 6150 - Advanced Topic Models
- MATH 7740 - Statistical Learning Theory
- CS 7790 - AI Seminar
- CS 7792 - Special Topics in Machine Learning